

Hw6

Cody

November 6, 2017

Stat 547M Homework 6:

Goal: Pick (at least) two of the six (numbered) topics below and do one of the exercise prompts listed, or something comparable using your dataset of choice.

```
library(tidyverse)
```

```
## Loading tidyverse: ggplot2
## Loading tidyverse: tibble
## Loading tidyverse: tidyr
## Loading tidyverse: readr
## Loading tidyverse: purrr
## Loading tidyverse: dplyr
```

```
## Conflicts with tidy packages -----
```

```
## filter(): dplyr, stats
## lag():    dplyr, stats
```

```
library(stringr)
```

1. Character data

Read and work the exercises in the Strings chapter or R for Data Science.

Exercises 1

1. In code that doesn't use stringr, you'll often see `paste()` and `paste0()`. What's the difference between the two functions? What stringr function are they equivalent to? How do the functions differ in their handling of NA?

A: `paste()` and `paste0()` have a different default value set for the 'sep' argument.

```
paste("hello", "world") #Default sep=" "
```

```
## [1] "hello world"
```

```
paste0("hello", "world") #Default sep=""
```

```
## [1] "helloworld"
```

The stringr function is `str_c()`

```
str_c("hello", "world")
```

```
## [1] "helloworld"
```

So we see that the default separator is also `sep=""`. If NAs are involved, `str_c()` will refuse to join where the others will treat them as strings. The NA has to be dealt with with a different function, `str_replace_na()`. This will turn an NA into a string "NA". Once this is done, then all functions behave similar.

```
paste("Hello",NA)
```

```
## [1] "Hello NA"
```

```
paste0("Hello",NA)
```

```
## [1] "HelloNA"
```

```
str_c("Hello",NA)
```

```
## [1] NA
```

```
paste("Hello",str_replace_na(NA))
```

```
## [1] "Hello NA"
```

```
paste0("Hello",str_replace_na(NA))
```

```
## [1] "HelloNA"
```

```
str_c("Hello",str_replace_na(NA))
```

```
## [1] "HelloNA"
```

2. In your own words, describe the difference between the `sep` and `collapse` arguments to `str_c()`.

`sep` will induce the type of spacing between the joined strings. For example,

```
str_c(c("This", "an"), c("is", "example"), sep="+")
```

```
## [1] "This+is"      "an+example"
```

But clearly we would have wanted this to be one string and this is where the `collapse` argument will put the separate strings together.

```
str_c(c("This", "an"), c("is", "example"), collapse="+")
```

```
## [1] "Thisis+anexample"
```

Together, we can make the correct string,

```
str_c(c("This", "an"), c("is", "example"), sep=" ", collapse=" ")
```

```
## [1] "This is an example"
```

3. Use `str_length()` and `str_sub()` to extract the middle character from a string. What will you do if the string has an even number of characters?

A: For the even length string, I'd return both values around the middle value.

```
exampleodd <- "Strings"
exampleeven <- "Cody"
```

```
odddlength <- str_length(exampleodd)
oddanswer <- ceiling(odddlength/2)
evenlength <- str_length(exampleeven)
evenanswer <- evenlength/2
```

```
str_sub(exampleodd, oddanswer, oddanswer)
```

```
## [1] "i"
```

```
str_sub(exampleeven,evenanswer,evenanswer+1)
```

```
## [1] "od"
```

4. What does `str_wrap()` do? When might you want to use it?

A: `str_wrap` is a way to break a string into smaller pieces but by following a special algorithm for keeping words **together** “Knuth-Plass wrapping”. This is useful for extracting information out of larger strings like a book or article for example.

5. What does `str_trim()` do? What’s the opposite of `str_trim()`?

A: This trims whitespace from strings.

```
string <- "          hi          "  
trimmed <- str_trim(string)  
padded <- str_pad(trimmed,width=10,side="both")
```

```
string
```

```
## [1] "          hi          "
```

```
trimmed
```

```
## [1] "hi"
```

```
padded
```

```
## [1] "    hi    "
```

6. Write a function that turns (e.g.) a vector `c("a", "b", "c")` into the string `a, b, and c`. Think carefully about what it should do if given a vector of length 0, 1, or 2.

A: This function will return an error for length 0, just the vector for length 1 and will combine following the above rule for any vector of size ≥ 2 .

```
my_func <- function(string_vec){  
  L <- length(string_vec)  
  if(L==0) {  
    stop('The length of the string is 0')  
  }  
  if(L==1){  
    return(string_vec)  
  }  
  firstpart <-  
    str_c(string_vec[1:L-1], collapse = ", ")  
  secondpart <-  
    str_c("and", string_vec[L], sep = " ")  
  final <-  
    str_c(firstpart, secondpart, sep = ", ")  
  return(final)  
}
```

To test it:

```
test <- letters[1:10]  
my_func(test)
```

```
## [1] "a, b, c, d, e, f, g, h, i, and j"
```

Exercises 2

1. Explain why each of these strings don't match a : `"", "\\", "\\", "`.

“" Doesn't work because a backslash is recognized as the beginning of a special character. This will look for the next character.”
 “\" Doesn't work because this now escapes from itself and creates a plain backslash, not a regular expression. “\" Now has the issue of forming a single pair to create a backslash but the unpaired backslash now is looking at the next character.

2. How would you match the sequence “’”?

```
x <- "\"'\\\"
writeLines(x)
```

" '\

3. What patterns will the regular expression match? How would you represent it as a string?

This will match anything that looks like “a.a.a”.

```
x <- "\\..\\.\\.\\.\\.\\.\\."
writeLines(x)
```

\. . \. . \. .

Exercises 3

1. How would you match the literal string “\$^\$”

```
x <- "$~$"  
writeLines(x)
```

$\$^{\wedge}\$$

```
str_match(x, "\\$\\|^\\$")
```

```
##      [,1]
## [1,] "$^$"
```

2. Create regular expressions that find:

- Start with “y”
- Ends with “x”
- Are exactly three letters long
- Have seven letters or more

```
test <- c("york","latex","abc","a lot of letters")
str_match(test,"^y")
```

```
##      [,1]
## [1,] "y"
## [2,] NA
## [3,] NA
## [4,] NA
```

```
str_match(test, "x$")
```

```
##      [,1]
## [1,] NA
## [2,] "x"
## [3,] NA
```

```
## [4,] NA
str_match(test, "^...$")
```

```
##      [,1]
## [1,] NA
## [2,] NA
## [3,] "abc"
## [4,] NA
```

```
str_match(test, ".....")
```

```
##      [,1]
## [1,] NA
## [2,] NA
## [3,] NA
## [4,] "a lot o"
```

Exercises 4

1. Find all words that contain:

- Start with a vowel
- Only consonants
- End with “ed” but not with eed
- End with “ing” or “ise”

```
test <- c("ao", "qwrt", "bed", "beed", "matching", "ise")
str_match(test, "^[aeiou]")
```

```
##      [,1]
## [1,] "a"
## [2,] NA
## [3,] NA
## [4,] NA
## [5,] NA
## [6,] "i"
```

```
str_match(test, str_c(rep("[^aeiou]", 4), collapse=""))
```

```
##      [,1]
## [1,] NA
## [2,] "qwrt"
## [3,] NA
## [4,] NA
## [5,] NA
## [6,] NA
```

```
str_match(test, "[^e][e][d]$")
```

```
##      [,1]
## [1,] NA
## [2,] NA
## [3,] "bed"
## [4,] NA
## [5,] NA
## [6,] NA
```

```
str_match(test,"([i][n][g])|([i][s][e])$")
```

```
##      [,1]  [,2]  [,3]
## [1,] NA    NA    NA
## [2,] NA    NA    NA
## [3,] NA    NA    NA
## [4,] NA    NA    NA
## [5,] "ing" "ing" NA
## [6,] "ise" NA    "ise"
```

Exercises 5

1. Describe the equivalences of “?”, “+” and “*” in {m,n} format.

A: ? - {0,1}, + - {1,}, * - {0,}

5. Work with a list

Here I will go through Trump Android Tweets

Loading the data

```
library(purrr)
suppressMessages(library(dplyr))
library(tibble)

load(url("http://varianceexplained.org/files/trump_tweets_df.rda"))
glimpse(trump_tweets_df)
```

```
## Observations: 1,512
## Variables: 16
## $ text          <chr> "My economic policy speech will be carried live ...
## $ favorited     <lgl> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE,...
## $ favoriteCount <dbl> 9214, 6981, 15724, 19837, 34051, 29831, 19223, 1...
## $ replyToSN     <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ...
## $ created       <dtm> 2016-08-08 15:20:44, 2016-08-08 13:28:20, 2016-...
## $ truncated     <lgl> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE,...
## $ replyToSID    <lgl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ...
## $ id           <chr> "762669882571980801", "762641595439190016", "762...
## $ replyToUID    <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ...
## $ statusSource  <chr> "<a href=\"http://twitter.com/download/android\""...
## $ screenName    <chr> "realDonaldTrump", "realDonaldTrump", "realDonal...
## $ retweetCount  <dbl> 3107, 2390, 6691, 6402, 11717, 9892, 5784, 7930,...
## $ isRetweet     <lgl> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE,...
## $ retweeted     <lgl> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE,...
## $ longitude     <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ...
## $ latitude      <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ...
```

```
tweets <- trump_tweets_df$text
tweets %>% head() %>% strtrim(70)
```

```
## [1] "My economic policy speech will be carried live at 12:15 P.M. Enjoy!"
## [2] "Join me in Fayetteville, North Carolina tomorrow evening at 6pm. Ticke"
## [3] "#ICYMI: \"Will Media Apologize to Trump?\" https://t.co/ia7rKBmioA"
```

```
## [4] "Michael Morell, the lightweight former Acting Director of C.I.A., and "
## [5] "The media is going crazy. They totally distort so many things on purpo"
## [6] "I see where Mayor Stephanie Rawlings-Blake of Baltimore is pushing Cro"

regex <- "badly|crazy|weak|spent|strong|dumb|joke|guns|funny|dead"
```

Here we select a few rows with the correct complexity to learn from

```
tweets <- tweets[c(1, 2, 5, 6, 198, 347, 919)]
tweets %>% strtrim(70)
```

```
## [1] "My economic policy speech will be carried live at 12:15 P.M. Enjoy!"
## [2] "Join me in Fayetteville, North Carolina tomorrow evening at 6pm. Ticke"
## [3] "The media is going crazy. They totally distort so many things on purpo"
## [4] "I see where Mayor Stephanie Rawlings-Blake of Baltimore is pushing Cro"
## [5] "Bernie Sanders started off strong, but with the selection of Kaine for"
## [6] "Crooked Hillary Clinton is unfit to serve as President of the U.S. Her"
## [7] "The Cruz-Kasich pact is under great strain. This joke of a deal is fal"
```

In base r

```
matches <- gregexpr(regex, tweets)
str(matches)
```

```
## List of 7
## $ : atomic [1:1] -1
##   .. attr(*, "match.length")= int -1
##   .. attr(*, "useBytes")= logi TRUE
## $ : atomic [1:1] -1
##   .. attr(*, "match.length")= int -1
##   .. attr(*, "useBytes")= logi TRUE
## $ : atomic [1:1] 20
##   .. attr(*, "match.length")= int 5
##   .. attr(*, "useBytes")= logi TRUE
## $ : atomic [1:1] 134
##   .. attr(*, "match.length")= int 4
##   .. attr(*, "useBytes")= logi TRUE
## $ : atomic [1:2] 28 95
##   .. attr(*, "match.length")= int [1:2] 6 4
##   .. attr(*, "useBytes")= logi TRUE
## $ : atomic [1:2] 87 114
##   .. attr(*, "match.length")= int [1:2] 4 6
##   .. attr(*, "useBytes")= logi TRUE
## $ : atomic [1:3] 50 112 123
##   .. attr(*, "match.length")= int [1:3] 4 4 4
##   .. attr(*, "useBytes")= logi TRUE
```

```
matches[[7]]
```

```
## [1] 50 112 123
## attr("match.length")
## [1] 4 4 4
## attr("useBytes")
## [1] TRUE
```

Base r isn't great. matches is quite ugly and hard to deal with.

```
lengths(matches) #happens to exist

## [1] 1 1 1 1 2 2 3
sapply(matches,length) #not friendly

## [1] 1 1 1 1 2 2 3
vapply(matches,length,integer(1)) #base approach

## [1] 1 1 1 1 2 2 3
map_int(matches,length) #purrr way

## [1] 1 1 1 1 2 2 3
```

Exercise: Get a list of the match lengths

```
match_length <-
  map(matches,~ attr(.x, which = "match.length"))
match_length

## [[1]]
## [1] -1
##
## [[2]]
## [1] -1
##
## [[3]]
## [1] 5
##
## [[4]]
## [1] 4
##
## [[5]]
## [1] 6 4
##
## [[6]]
## [1] 4 6
##
## [[7]]
## [1] 4 4 4
```

Exercise: Count “Trump words”

Write a function that will count the number of these “Trump words”.

```
f <- function(x) sum(x>0)
map(matches,f)

## [[1]]
## [1] 0
##
## [[2]]
## [1] 0
```



```
##
## [[3]]
## [1] 1
##
## [[4]]
## [1] 1
##
## [[5]]
## [1] 2
##
## [[6]]
## [1] 2
##
## [[7]]
## [1] 3
```

But a more compact way to get a vector of integers, we can use the purrr function:

```
map_int(matches, ~sum(.x>0))
```

```
## [1] 0 0 1 1 2 2 3
```

Excercise: Strip the attributes from matches

Here we use the as.vector() function which will throw away non-vectorizable structure

```
match_first <- map(matches, as.vector)
str(match_first)
```

```
## List of 7
## $ : int -1
## $ : int -1
## $ : int 20
## $ : int 134
## $ : int [1:2] 28 95
## $ : int [1:2] 87 114
## $ : int [1:3] 50 112 123
```

Exercise: Find the final word that fits the “Trump word” criteria

```
match_last <- map2(match_first, match_length, ~ .x+.y-1)
match_last
```

```
## [[1]]
## [1] -3
##
## [[2]]
## [1] -3
##
## [[3]]
## [1] 24
##
## [[4]]
## [1] 137
```

```
##
## [[5]]
## [1] 33 98
##
## [[6]]
## [1] 90 119
##
## [[7]]
## [1] 53 115 126
```

Put it all together using pmap

This will find all the “Trump words” in a tweet and pull them out.

```
pmap(list(text=tweets,first=match_first,last=match_last),.f=substring)
```

```
## [[1]]
## [1] ""
##
## [[2]]
## [1] ""
##
## [[3]]
## [1] "crazy"
##
## [[4]]
## [1] "joke"
##
## [[5]]
## [1] "strong" "weak"
##
## [[6]]
## [1] "weak" "strong"
##
## [[7]]
## [1] "joke" "dead" "dumb"
```

Just to put it all in one place,

```
trump_frame <-tibble(
  text = tweets,
  first = match_first,
  last = match_last
)
pmap(trump_frame, substring)
```

```
## [[1]]
## [1] ""
##
## [[2]]
## [1] ""
##
## [[3]]
## [1] "crazy"
##
```

```
## [[4]]
## [1] "joke"
##
## [[5]]
## [1] "strong" "weak"
##
## [[6]]
## [1] "weak"   "strong"
##
## [[7]]
## [1] "joke" "dead" "dumb"
```

Reflection

! Regular expressions are rough! I've learned them now two or three times and they're never fun! I didn't work through every single example as there are so many but I tried to do at least one in each section. I still probably need a formal lecture on some of the specifics, like grouping and backreferencing.

As for using purrr functions over strings, I found this to be incredibly easy and I can see why there are better alternatives to base R.